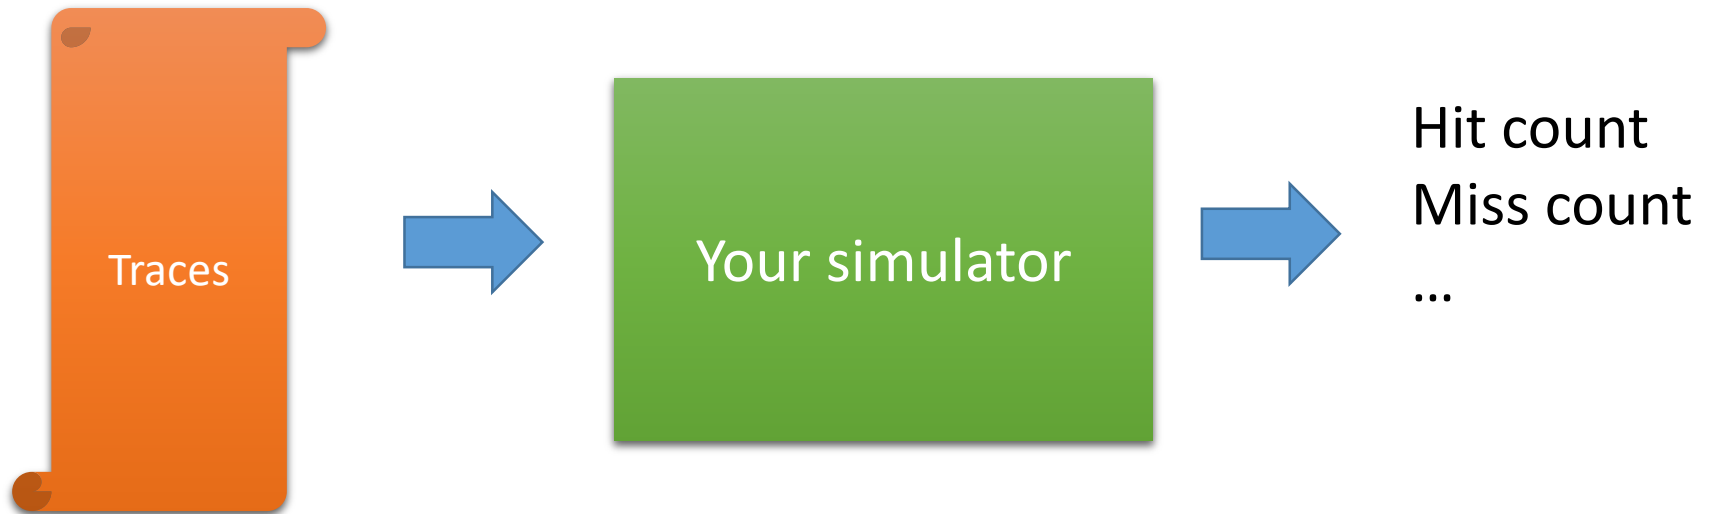


Operating Systems Programming Assignment #5

Page Replacement Simulation: LRU and LFU

Prof. Li-Pin Chang, NCTU


Simulation



Trace File Format

- The trace format and Zipfian distribution

1003
1003
9340
1243
1108
1786
1066
1312
1000
1000
1213
1249
2116



Unsigned integer

Page Replacement(LFU)

- Example: Frame #=2

4001 (miss)

LFU	
4001	

Ref count		
4000	4001	4002
0	1	0

4001 (hit)

4001	
------	--

0	2	0
---	---	---

4000 (miss)

4000	4001
------	------

1	2	0
---	---	---

4002 (miss)

4000 ←	4002	4001
--------	------	------


0	2	1
---	---	---

4000 (miss)

4002 ←	4000	4001
--------	------	------

1	2	0
---	---	---

Simulator Structure (LFU)



Page Replacement(LRU)

- Example: Frame #=2


4001 (miss)

4001 (hit)


4000 (miss)

4002 (miss)

4000 (hit)



Simulator Structure (LRU)



Page Cache Operations

- Page lookup
 - Check whether or a new reference is a hit or a miss
 - Hash tables, binary search trees, skip lists....
- Do not use linear search!!!
 - You will receive a grade penalty if you do
 - Implement your own search, or reuse any existing libraries/classes for searching
 - TAs will examine your code
 - Duplication in this part does not count

Victim selection

- LFU
 - The least frequently used page
 - If two pages have the same access count, the page having a smaller reference sequence number is replaced
 - You may need to store the reference sequence number when a page is added to the page cache
- LRU
 - The least recently used page

Procedure

1. Algorithm=LFU
2. For (frame # = 128, 256, 512, and 1024) do
 - Read the trace file
 - Run simulation
 - Print out the hit count, miss count, page fault ratio
3. Print out the total elapsed time of Step 2
4. Algorithm=LRU
5. For (frame # = 128, 256, 512, and 1024) do
 - Read the trace file
 - Run simulation
 - Print out the hit count, miss count, page fault ratio
6. Print out the total elapsed time of Step 5

The scenario of program

Output format

LFU policy:(\n)

frame	(\t)	hit	(\t\t)	miss	(\t\t)	page fault ratio	(\n)
128	(\t)	%d	(\t\t)	%d	(\t\t)	%.10f	(\n)
256	(\t)	%d	(\t\t)	%d	(\t\t)	%.10f	(\n)
512	(\t)	%d	(\t\t)	%d	(\t\t)	%.10f	(\n)
1024	(\t)	%d	(\t\t)	%d	(\t\t)	%.10f	(\n)
Total elapsed time: %.4f sec (\n)							

LRU policy:(\n)

frame	(\t)	hit	(\t\t)	miss	(\t\t)	page fault ratio	(\n)
128	(\t)	%d	(\t\t)	%d	(\t\t)	%.10f	(\n)
256	(\t)	%d	(\t\t)	%d	(\t\t)	%.10f	(\n)
512	(\t)	%d	(\t\t)	%d	(\t\t)	%.10f	(\n)
1024	(\t)	%d	(\t\t)	%d	(\t\t)	%.10f	(\n)
Total elapsed time: %.4f sec (\n)							

The scenario of program

- Assume the path of the trace file is `./test/sample.txt`
Your output should look like this

```
[tsaichh0619@linux1 assignment5]$ ./bin/ta_5 ./test/sample.txt
LFU policy:
frame    hit          miss          page fault ratio
128      8382341          1617659          0.1617659032
256      8807904          1192096          0.1192096025
512      9134228          865772           0.0865771994
1024     9404369          595631           0.0595631003
Total elapsed time: 2.4703 sec

LRU policy:
frame    hit          miss          page fault ratio
128      7880806          2119194          0.2119193971
256      8425225          1574775          0.1574774981
512      8871064          1128936          0.1128935963
1024     9236112          763888           0.0763887987
Total elapsed time: 2.1897 sec
```

Correctness

- Except the total elapsed time, your results should be exactly the same as ours
- Do not use linear search in anywhere of your program; otherwise, you will receive a score penalty
- Remark: here we do not simulate working set migration, so LFU slightly outperforms LRU; but in real life workloads, LRU usually slightly outperforms LFU

More details

- Total request # \leq 500 millions
- Total frame # \leq 8192
- The path+file name of the trace file is an argument of your program (see the screen shot), do not hard-coding the pathname of the trace file
- For each iteration, you should open the file, run the simulation, print the result and close the file
- Do not store the trace data to speed up the next iteration
- Use `gettimeofday()` to get the total elapsed time

Header of your .c or .cpp

```
/*
```

```
Student No.: <your student id>
```

```
Student Name: <your name>
```

```
Email: <your email>
```

```
SE tag: xnxctxuxoxsx
```

```
Statement: I am fully aware that this program is not  
supposed to be posted to a public server, such as a  
public GitHub repository or a public web page.
```

```
*/
```

About Submission

- We will compile your code on **Ubuntu 16.04** or **CS linux work station**
- You should name your program into **xxx_5.c** or **xxx_5.cpp** (**xxx** is your student ID)
- How we compile your program:
“ gcc -std=gnu99 xxx_5.c -o xxx_5 ” or
“ g++ -std=c++11xxx_5.cpp -o xxx_5 ”
- If you use a make file, please upload **one single zip file** that contains your make file and source file (the same naming convention as above)

Plagiarism

- We take plagiarism seriously
- You will get zero points for plagiarism
- “Copying existing code from the internet and revising it for submission” is plagiarism
- “Consulting to existing code from the internet and writing a new program from scratch” is not